

Database Systems

Concurrency

H. Turgut Uyar Şule Öğüdücü

2002-2012

1 / 45

License



©2002-2012 T. Uyar, Ş. Öğüdücü

You are free:

- ▶ to Share – to copy, distribute and transmit the work
- ▶ to Remix – to adapt the work

Under the following conditions:

- ▶ Attribution – You must attribute the work in the manner specified by the author or licensor (but not in any way that suggests that they endorse you or your use of the work).
- ▶ Noncommercial – You may not use this work for commercial purposes.
- ▶ Share Alike – If you alter, transform, or build upon this work, you may distribute the resulting work only under the same or similar license to this one.

Legal code (the full license):

<http://creativecommons.org/licenses/by-nc-sa/3.0/>

2 / 45

Topics

Transactions

Introduction
Recovery
Two-Phase Commit

Concurrency

Introduction
Locking
Isolation Levels
Intent Locks

3 / 45

Transactions

- ▶ a group of operations may have to be carried out together
 - ▶ finishing some operations while failing on others might cause inconsistency
- ▶ no guarantee that multiple operations will all be finished
 - ▶ at least return to the state before the changes

Definition

transaction: a logical unit of work

4 / 45

Transaction Example

Example (transferring money from one bank account to another)

```
UPDATE ACCOUNTS SET BALANCE = BALANCE - 100
WHERE ACCOUNTID = 123
```

```
UPDATE ACCOUNTS SET BALANCE = BALANCE + 100
WHERE ACCOUNTID = 456
```

5 / 45

Transaction Properties

- ▶ atomicity
 - ▶ all or nothing
- ▶ consistency
 - ▶ moving from one consistent state to another
- ▶ isolation
 - ▶ whether operations of an unfinished transaction affect other transactions or not
- ▶ durability
 - ▶ when a transaction is finished, its changes are permanent even if there is a system failure

6 / 45

Transaction Start/End

starting a transaction

```
BEGIN [ WORK | TRANSACTION ]
```

finishing a transaction

```
COMMIT [ WORK | TRANSACTION ]
```

cancelling a transaction

```
ROLLBACK [ WORK | TRANSACTION ]
```

7 / 45

Transaction Example

Example

```
BEGIN TRANSACTION
ON ERROR GOTO UNDO
UPDATE ACCOUNTS SET BALANCE = BALANCE - 100
  WHERE (ACCOUNTID = 123)
UPDATE ACCOUNTS SET BALANCE = BALANCE + 100
  WHERE (ACCOUNTID = 456)
COMMIT
...
```

```
UNDO :
ROLLBACK
```

8 / 45

Recovery

- ▶ consider a system failure during a transaction
 - ▶ buffer cache has not been flushed to the disk
- ▶ how to guarantee durability?
 - ▶ it should be possible to derive the data from other sources in the system
 - ▶ internal level

9 / 45

Transaction Log

- ▶ the **log** keeps the values of every affected tuple before and after every operation
- ▶ **write-ahead log rule**:
the log must be flushed to the physical medium before the transaction is committed
- ▶ accessing records in the log is sequential by nature

10 / 45

Checkpoints

- ▶ create **checkpoints** in the log at certain intervals
- ▶ flush buffer cache to the physical medium
- ▶ note the checkpoint in the log
- ▶ note the continuing transactions

11 / 45


Recovery Lists

- ▶ after the failure, which transactions will be undone, which transactions will be made permanent?
 - ▶ create two lists: *undo* (U), *redo* (R)
- ▶ t_C : last checkpoint in the log
 - ▶ add the transactions which are active at t_C to U
- ▶ scan records from t_C to end of log
 - ▶ add any starting transaction to U
 - ▶ move any finishing transaction to R

12 / 45

Recovery Example

Example



- ▶ t_C :
 $U = \{T_2, T_3\}$ $R = \emptyset$
- ▶ T_4 started:
 $U = \{T_2, T_3, T_4\}$ $R = \emptyset$
- ▶ T_2 finished:
 $U = \{T_3, T_4\}$ $R = \{T_2\}$
- ▶ T_5 started:
 $U = \{T_3, T_4, T_5\}$ $R = \{T_2\}$
- ▶ T_4 finished:
 $U = \{T_3, T_5\}$ $R = \{T_2, T_4\}$

13 / 45

Recovery Process

- ▶ scan records backwards from end of log
 - ▶ undo the changes made by the transactions in U
- ▶ scan records forwards
 - ▶ redo the changes made by the transactions in R

14 / 45

Two-Phase Commit

- ▶ different source managers
 - ▶ different undo / redo mechanisms
- ▶ modifications on data that reside on different source managers
 - ▶ either commit in all source managers or rollback in all source managers
- ▶ coordinator

15 / 45

Protocol

- ▶ coordinator tells all participants to flush all data regarding the transaction to the physical medium
- ▶ coordinator tells all participants to start the transaction and report back the result
 - ▶ if all participants report success, coordinator decides to commit
 - ▶ if one or more participants report failure, coordinator decides to rollback
- ▶ coordinator informs the participants about the decision

16 / 45

References

Required Reading: Date

- ▶ Chapter 15: Recovery

17 / 45

Concurrency

- ▶ problems that might arise due to simultaneous transactions:
 - ▶ lost update
 - ▶ uncommitted dependency
 - ▶ inconsistent analysis

18 / 45

Lost Update

Example

Transaction A	Transaction B
...	...
RETRIEVE p	...
...	...
...	RETRIEVE p
...	...
UPDATE p	...
...	...
...	UPDATE p
...	...

19 / 45

Uncommitted Dependency

Example

Transaction A	Transaction B
...	...
...	UPDATE p
...	...
RETRIEVE p	...
...	...
...	ROLLBACK
...	...

20 / 45

Inconsistent Analysis

Example (sum of accounts: acc1=40, acc2=50, acc3=30)

Transaction A	Transaction B
...	...
RETRIEVE acc1 (40)	...
RETRIEVE acc2 (90)	...
...	...
...	UPDATE acc3 (30 → 20)
...	UPDATE acc1 (40 → 50)
...	COMMIT
...	...
RETRIEVE acc3 (110)	...
...	...

21 / 45

Conflicts

- ▶ A reads, B reads
 - ▶ no problem
- ▶ A reads, B writes
 - ▶ non-repeatable read (inconsistent analysis)
- ▶ A writes, B reads
 - ▶ dirty read (uncommitted dependency)
- ▶ A writes, B writes
 - ▶ dirty write (lost update)

22 / 45

Locking

- ▶ transactions lock the tuples they work on
 - ▶ shared lock (S)
 - ▶ exclusive lock (X)
- ▶ they release the locks when they are done

23 / 45

Lock Requests

lock type compatibility matrix

	X	S	-
X	N	N	Y
S	N	Y	Y

- ▶ if exclusive lock, all lock requests are denied
- ▶ if shared lock:
 - ▶ exclusive lock requests are denied
 - ▶ shared lock requests are granted

24 / 45

Locking Protocol

- ▶ the transaction requests a lock depending on the operation it wants to perform
 - ▶ promote a shared lock to an exclusive lock
- ▶ if the request cannot be granted, it starts waiting
 - ▶ it continues when the transaction that holds the lock releases it
 - ▶ **starvation**

25 / 45

Lost Update

Example

Transaction A	Transaction B
...	...
RETRIEVE p (S+)	...
...	...
...	RETRIEVE p (S+)
...	...
UPDATE p (X-)	...
wait	...
wait	UPDATE p (X-)
wait	wait

26 / 45

Uncommitted Dependency

Example

Transaction A	Transaction B
...	...
...	UPDATE p (X+)
...	...
RETRIEVE p (S-)	...
wait	...
wait	ROLLBACK
RETRIEVE p (S+)	...
...	...

27 / 45

Inconsistent Analysis

Example (sum of accounts: acc1=40, acc2=50, acc3=30)

Transaction A	Transaction B
...	...
RETRIEVE acc1 (S+)	...
RETRIEVE acc2 (S+)	...
...	...
...	UPDATE acc3 (X+)
...	UPDATE acc1 (X-)
...	wait
RETRIEVE acc3 (S-)	wait
wait	wait

28 / 45

Deadlock

Definition

deadlock:


transactions waiting for each other to release the locks

- ▶ almost always between two transactions
- ▶ countermeasures:
 - ▶ detecting and solving
 - ▶ preventing

29 / 45

Solving Deadlocks

Example



- ▶ wait graph
- ▶ choose a **victim** and kill it

30 / 45

Preventing Deadlocks

- ▶ every transaction has a starting timestamp
- ▶ if the lock request of transaction A conflicts with a lock held by transaction B:
 - ▶ **wait-die**: A waits if it is older than B, otherwise it dies
A is rolled back and restarted
 - ▶ **wound-wait**: A waits if it is younger than B, otherwise it wounds B
B is rolled back and restarted
- ▶ the timestamp of a restarted transaction is not changed

31 / 45

Lock Statements

shared lock

```
SELECT query FOR SHARE
```

exclusive lock

```
SELECT query FOR UPDATE
```

32 / 45

Isolation Levels

- ▶ if isolation is decreased, concurrency can be increased
- ▶ various isolation levels:
 - ▶ serializable
 - ▶ repeatable read
 - ▶ read committed
 - ▶ read uncommitted

33 / 45

Serializability

- ▶ **serial execution**:
a transaction starts only after another is finished
- ▶ **serializable**: the result of concurrent execution is always the same as one of the serial executions

Example

- ▶ $x = 10$
- ▶ transaction A: $x = x + 1$
- ▶ transaction B: $x = 2 * x$
- ▶ first A, then B: $x = 22$
- ▶ first B, then A: $x = 21$

34 / 45

Two-Phase Locking

- ▶ **two-phase locking**:
after any lock is released
there will be no more new lock requests
 - ▶ expansion phase: gather locks
 - ▶ contraction phase: release locks
- ▶ **two-phase strict locking**:
all locks are released at the end of the transaction
- ▶ *If all transactions obey the two-phase locking protocol, all concurrent executions are serializable.*

35 / 45

Read Committed

- ▶ only exclusive locks are held until end of the transaction

Example

Transaction A	Transaction B
...	...
RETRIEVE p (S+)	...
...	...
release lock	...
...	...
...	UPDATE p (X+)
...	COMMIT
RETRIEVE p (S+)	

36 / 45

Phantoms

Definition

phantom: when the query is executed again, new tuples appear

Example

- ▶ transaction A computes the average of a customer's account balances:
 $\frac{100+100+100}{3} = 100$
- ▶ transaction B creates new account with balance 200 for the same customer
- ▶ transaction A computes again:
 $\frac{100+100+100+200}{4} = 125$

37 / 45

Setting Isolation Levels

Statement

```
SET TRANSACTION ISOLATION LEVEL
[ SERIALIZABLE | REPEATABLE READ |
  READ COMMITTED | READ UNCOMMITTED ]
```

38 / 45

Isolation Level Problems

isolation level	dirty read	non-repeatable read	phantom
READ UNCOMMITTED	Y	Y	Y
READ COMMITTED	N	Y	Y
REPEATABLE READ	N	N	Y
SERIALIZABLE	N	N	N

39 / 45

Locking Granularity

- ▶ locking relations instead of tuples
 - ▶ even the entire database
- ▶ if granularity is increased, concurrency is decreased
- ▶ hard to find locks on tuples
 - first, get **intent locks** on relation variables

40 / 45

Intent Locks

- ▶ Intent Shared (IS):
the transaction intends to read some tuples
- ▶ Intent Exclusive (IX):
IS + the transaction intends to write some tuples
- ▶ Shared (S):
concurrent readers are allowed but no concurrent writers
- ▶ Shared + Intent Exclusive (SIX):
S + IX
- ▶ Exclusive (X):
no concurrency allowed on this relation

41 / 45


Lock Requests

lock compatibility matrix

	X	SIX	IX	S	IS	-
X	N	N	N	N	N	Y
SIX	N	N	N	N	Y	Y
IX	N	N	Y	N	Y	Y
S	N	N	N	Y	Y	Y
IS	N	Y	Y	Y	Y	Y

42 / 45

Lock Precedence



- ▶ for a shared lock on a tuple, at least an IS lock on the relation
- ▶ for an exclusive lock on a tuple, at least an IX lock on the relation

43 / 45

Locking Statements

Statement

```
LOCK [ TABLE ] table_name  
[ IN lock_mode MODE ]
```

- ▶ lock modes:
 - ▶ ACCESS SHARE
 - ▶ ROW SHARE
 - ▶ ROW EXCLUSIVE
 - ▶ SHARE UPDATE EXCLUSIVE
 - ▶ SHARE
 - ▶ SHARE ROW EXCLUSIVE
 - ▶ EXCLUSIVE
 - ▶ ACCESS EXCLUSIVE

44 / 45

References

Required Reading: Date

- ▶ Chapter 16: Concurrency

45 / 45